



Build with Gemini Hackathon 2026 — Participant Guide

Your single reference for the day: event schedule, rules, prize tracks, model guidance, judging, and submission requirements.

⚡ Build something extraordinary with Gemini and Gemma.

[Google DeepMind](#) × [65labs](#)

Date: @Today

Venue: Lorong AI @ One-North, Singapore

Submission deadline: 3:30 pm sharp

► [Jump to a section](#)

The essentials

What	Details
Team size	Up to four members
Build requirement	Start the project during today's hackathon—new work only
Submission deadline	3:30 pm; no late entries
Prize pool	US\$3,000 total: one US\$1,000 winner in each of three tracks
Required deliverables	Public GitHub repository, short demo video (maximum 3:00 / three minutes), project name, one-line description, and selected track(s)—at most two
Core technology rule	Gemini or Gemma must be essential to what the product does



Google AI Studio credits: Watch your email for the personalized username and password for your temporary developer account. Sign in to [Google AI Studio](#) with those credentials, then **create an API key in AI Studio**. Your **US\$25 in credits is already configured**—there is no need to visit Google Cloud Console/GCP or create or configure a Cloud project. Please keep your login details private.



Guest Wi-Fi: Vidacity-IorongAI Guest Wifi

Password: 4xgtnk

Participant communications



QR code to join the participant
WhatsApp group



Join the participant WhatsApp group

Scan the QR code for event updates,
schedule changes, and participant
communications throughout the day.

[Open the WhatsApp group](#)

Today's schedule

Time	What's happening
8:30 am	Registration, check-in, and refreshments
9:00 am	Kickoff, rules, tracks, and prizes
9:10 am	Hacking begins
12:00 pm	Lunch buffet opens. Take food when you're ready; keep building if you prefer.
3:30 pm	Submissions close
3:30–4:30 pm	Preliminary judging
4:30–5:30 pm	Finalists announced and presentations
5:30–6:00 pm	Awards and departure

Coffee and extra workspace

Lume Coffee, one floor down, is open from **10:00 am–4:00 pm** during the main hacking window.

- You're welcome to use it as additional workspace.
- Limited free coffee is available first come, first served.
- Drinks are also available for purchase.

Three tracks, three winners

Each track awards one **US\$1,000 prize**. Each project may enter **at most two tracks**.

Track	Core challenge	What your demo should prove
Best Use of Gemma	Push an open model somewhere thoughtful, useful, or technically surprising.	Gemma’s openness, deployability, privacy, efficiency, or adaptability is essential—not interchangeable.
Best Elderly Hack	Make modern technology easier for older adults to use, trust, and stay connected through.	A senior, caregiver, or eldercare worker can complete a meaningful task with low friction and appropriate safety.
Most Creative Gemini Hack	Build the wildest, fastest, or most original use of Gemini Flash 3.7.	Gemini’s speed or multimodal behavior creates a memorable live moment.

Track 1 — Best Use of Gemma

The brief

Open models, pushed further. Gemma has crossed one billion cumulative downloads. Show what an open model can do when you make a deliberate technical or product choice—not merely swap it into a generic demo.

Strong directions from the event brief:

- **Run privately or at the edge**
- **Adapt the model for a specific community**
- **Make openness essential to the product**

What Gemma gives you

The technical deck presents Gemma as an open-model family designed around efficiency and flexible deployment.

- **Multiple model sizes:** Choose a model that matches the device, latency, and quality constraints of the product.
- **Multimodal capability:** The deck highlights vision tasks such as object detection, pointing and labeling, and reading text in images.
- **Quantization-aware training:** QAT versions are designed to reduce VRAM and storage needs while preserving capabilities.
- **Multi-token prediction:** The deck describes an inference-speed optimization intended to reduce wait time, particularly for domains such as coding.
- **Flexible deployment:** Gemma can be used on mobile, in a web app, on a local machine, or in cloud/server environments including Firebase AI Logic, Cloud Run, and GKE.

The source deck discusses multiple Gemma generations and sizes. Confirm the exact models, weights, SDKs, licenses, and hardware available today before locking the architecture.

Strong project shapes

- **Private by design:** Sensitive documents, images, or audio stay on-device or in controlled infrastructure.
- **Works under constraint:** The experience runs with limited connectivity, memory, compute, latency, or cost.
- **Adapted to context:** The model serves a specific language, community, device, environment, or specialist workflow.
- **Technically transparent:** The demo exposes a meaningful choice—model size, quantization, response time, quality, or deployment target.

Idea sparks

- An offline visual guide that reads labels or detects hazards without sending images to the cloud.
- A private document simplifier for sensitive letters, with uncertainty and source passages shown.
- A field technician assistant built around one equipment manual for low-connectivity sites.
- A local-language community kiosk for one narrow civic or neighborhood workflow.

- A model-efficiency lab that recommends the smallest Gemma configuration meeting a real latency and quality target.

Make the Gemma choice visible

Your demo should answer: **Why Gemma?**

Good proof includes:

- Disconnecting the network and completing the task.
- Showing what data stays local.
- Running on the target device.
- Comparing two model or quantization choices on a tiny repeatable test.
- Demonstrating a community or domain adaptation that changes the product experience.

Track 2 — Best Elderly Hack

The opportunity

By 2030, one in six people globally will be aged 60+. This track asks teams to build empathetic, voice-first, and context-aware tools for seniors, family caregivers, and eldercare communities.

Three challenge areas

1. Voice-first and multimodal accessibility

Standard interfaces can be cluttered and text-heavy and may not handle slower speech, conversational pauses, or visual impairments well.

Build direction: Create a hands-free, patient interaction using voice, camera, or multimodal input.

Examples from the track brief:

- An ambient Gemini Live companion that assists with daily tasks without requiring app navigation.
- A camera-based visual aid that reads medication labels or official letters and explains them in plain speech.

2. Context-aware caregiving

Generic AI can give tone-deaf or unsafe advice.

Build direction: Use context carefully to deliver bounded, personalized, situation-aware support.

Examples from the track brief:

- A caregiver co-pilot that tracks changes in mood, meals, or medication routines and prepares a transparent summary.
- A dementia-friendly memory bridge that uses family photos and audio to support gentle conversation and preserve stories.

3. Family and community connection

Social isolation is a major risk for people aging alone.

Build direction: Foster cognitive vitality, intergenerational connection, and emotional wellbeing.

Examples from the track brief:

- An oral-history storyteller that helps grandparents turn spoken memories into illustrated, editable stories for their families.
- A daily wellness and trivia companion that can identify distress or inactivity and involve a trusted contact appropriately.

Official Elderly-track judging criteria

Criterion	Weight	Question to answer
Empathy and usability	40%	Is the interface genuinely accessible for a 70+ year old through low cognitive friction, voice-first interaction, high contrast, and clear language?
Contextual safety and reliability	30%	Does the model handle caregiving edge cases safely and avoid hallucination?
Real-world impact and feasibility	30%	Does it solve a tangible pain point for seniors, caregivers, or eldercare workers?

Safety and accessibility baseline



Health-related products may require appropriate regulatory approvals before being offered to users. For today's prototype, avoid diagnosis, medication changes, emergency claims, and unsupervised consequential actions. Show uncertainty, ask for confirmation, and route high-stakes decisions to a qualified human.

Before submitting, test the core flow with someone unfamiliar with it:

- Can they begin without reading a long explanation?
- Does it tolerate pauses, corrections, and incomplete speech?
- Is there one clear next action at a time?
- Can the user repeat, undo, or ask why?
- Does it distinguish known facts, model inference, and uncertainty?
- Is consent explicit before sharing information with family, caregivers, or professionals?

Track 3 — Most Creative Gemini Hack

The brief

The unexpected use of Gemini Flash 3.7. Surprise the room. Use Gemini's speed and multimodal capabilities to create an interaction, tool, or experience people will remember after the demos end.

Strong directions from the event brief:

- **Make the model behavior visible**
- **Create a memorable live moment**
- **Turn speed into a product advantage**

What Gemini gives you

The technical deck shows Gemini accepting a wide range of inputs:

- System instructions
- Text and code
- Images
- Audio and speech
- Video
- API information
- URLs

It also shows Gemini producing or triggering:

- Text, images, audio, and video
- Function calls
- Code

The deck positions Gemini Flash 3.7 around coding, production code quality, knowledge work, workflow automation, and fast multimodal interaction.

Strong project shapes

- The model sees or hears something live, reasons about it, and responds in the right medium.
- Mixed documents, URLs, images, or recordings become one coherent and traceable outcome.

- A function call completes an action rather than merely explaining what to do.
- Speed enables an interaction that would feel broken with a slow model.
- The audience can see the model adapt to an unexpected input.

Idea sparks

- Turn the objects in a real room into a playable collaborative game.
- Create a live improv conductor that reacts to dialogue and gestures with cues.
- Transform clips, photos, voice notes, and URLs into an editable mini-documentary.
- Build a promptless maker assistant that intervenes only when it detects confusion or risk.
- Generate a temporary interface after Gemini sees a physical setup or workflow.
- Fuse several people's sketches, sounds, gestures, or sentences into one evolving audiovisual world.

Design the “wow” moment first

A creative project does not need many features. It needs one strong interaction:

1. The audience clearly sees or hears the live input.
2. Gemini performs behavior that could not be reduced to a fixed script.
3. The response arrives quickly enough to feel interactive.
4. The output changes the experience, not merely the copy.

5. The team can explain what was live, cached, mocked, or precomputed.

What not to build

⊗ The event deck identifies the following categories as overdone or problematic and says to strictly avoid them:

- AI mental health advisor
- Basic RAG application
- Streamlit application
- Dashboard
- Image analyzer
- Chatbot, including “AI for Education”
- AI job application screener
- AI nutrition coach
- Personality analyzer

This does not mean that every interface containing retrieval, chat, images, or a dashboard component is automatically disqualified. It means those familiar patterns should not be the project’s core idea. **Be creative. Surprise the room.**

Judging & submission

Rules and fairness

1. Teams may have up to four members.
2. Start the project during today's hackathon—new work only.
3. Use only code, data, and assets that you have the right to use.
4. Projects that violate legal, ethical, or platform policies may be disqualified.
5. Each project may enter at most two tracks.
6. Submit a public GitHub repository and a short demo video (maximum 3:00 / three minutes) by 3:30 pm.

Preliminary judging with Gavel

During preliminary judging, judges will use **Gavel**, a project-expo judging system originally built for HackMIT. Instead of asking every judge to score every project, Gavel uses **pairwise comparisons**: judges compare projects against one another, and the system combines those decisions into a live ranking by inferred project quality.

This lets multiple judges assess a large field in parallel and helps us identify a finalist shortlist efficiently. The pairwise result supports the preliminary selection; finalists will then present during the final round. [Learn more about Gavel.](#)

What judges are looking for

- **Clear problem:** The audience, need, and intended payoff are easy to understand.
- **Working core loop:** Show the main experience first. Extra features can wait.
- **Meaningful AI use:** Gemini or Gemma is essential to what the product does.
- **Memorable demo:** Make one moment work well enough that the judges remember it.

The four judging lenses

Lens	What it means
Technical merit	Quality of code, architecture, and implementation
Polish	Attention to detail, UX, and overall execution
Execution	How well the team delivered on its vision
Wow factor	How memorable and impressive the hack is

Submission requirements



Hard deadline: 3:30 pm. No late entries. The submission link will be shared in the room.

Submit:

- A public GitHub repository
- A short demo video—maximum 3:00 (three minutes)
- Your **project name**
- A **one-line description**
- Your **selected track(s)**—at most two

Demo video examples

These are examples of concise product demos—not templates you need to copy. Keep your own video under 3:00, show the working code, be early, and make the model's role easy to