

Gemini Hackathon 2026: Official Selection Panel Evaluation Report

Singapore, 19 August 2026 — As the official, evidence-led selection panel for the **Singapore Build with Gemini Hackathon 2026** (co-hosted by 65labs, Google DeepMind DX, and Lorong AI) 1, we have completed a rigorous evaluation of all nine candidate ideas proposed across research-kimi.md, research-zai.md, and research-chatgpt.md.

To maintain absolute credibility and align with the hackathon's core principles, we have evaluated each proposal against the independent **Forensic Field Analysis of Budget Travel Pain Points** 2 and its referenced original sources in this notebook. Any claims originating from the candidate reports themselves have been treated as unverified unless validated by the notebook's independent source materials.

1. Candidate Inventory

Below is the inventory of all nine evaluated proposals:

#,Idea Name,Origin,Claimed Theme,One-Sentence User Outcome

1,TrueFare Lens 3,ChatGPT,Best Use Of Gemma 4,Travelers compare realistic all-in flight costs (including checked/carry-on bags and required extras) instead of misleading teaser fares 3.

2,Miss-It-Not 5,ChatGPT,Most Creative Gemini Hack 6,Travelers receive a quantified savings-to-risk audit for self-transfer itineraries to decide if tight layovers are worth the booking hazard 5.

3,Senior Fare Copilot 7,ChatGPT,Best Elderly Hack 8,Older travelers use a simplified voice interface to search and mock-book cheaper flights without navigating complex digital application sprawl 7.

4,GateMate 9,Kimi,Best Elderly Hack 10,"Elderly travelers capture a boarding pass photo, speak in their native dialect, and receive simplified, spoken gate-navigation guidance 9."

5,ClaimMate 11,Kimi,Best Use of Gemma 12,Delayed travelers snap photos of disruption emails and boarding passes to privately verify eligibility and draft statutory claim letters 11.

6,StitchCheck 13,Kimi,Most Creative Gemini Hack 14,"Travelers upload separate ticket screenshots to visualize gate-to-gate connection risks, map failure timelines, and secure protected alternatives 13."

7,TrueFare 15,Z.ai,Everyday AI 15,"Travelers view an itemized, all-in budget fare with every bag, seat, and check-in fee cited directly to carrier rules 15."

8,ConnectionGuard 16,Z.ai,The Next Frontier 16,"Travelers stress-test separate-ticket connections for terminal changes and delays, seamlessly securing safer alternative flights 16."

9,DoorFare 17,Z.ai,Social Good 18,"Travelers compare flights by true door-to-door cost, exposing hidden public transit shuttle fees and travel times to secondary airports 17."

2. Evidence Audit

We audited each idea's underlying passenger pain point against the notebook's independent sources to separate empirical evidence from model-generated assumptions:

- **TrueFare Lens (ChatGPT Candidate 1)**

- **Pain-Point Evidence Strength: Strong** (general pain point) / **Unsupported** (specific flight data) 19, 20.
- **Supporting Source Citations:** *Forensic Field Analysis* (Finding 1) 21, 22, *BEUC Factsheet* ("Fly light, pay heavy") 23, 24, *Which? Gate Bonus Report* 25.
- **Independent of Proposing Model?** Yes. General baggage fee drip pricing is thoroughly documented by the independent consumer body BEUC 23 and the *Forensic Field Analysis* 21.
- **Audit Gaps & Contradictions:** The report claims a "Which? December 2025 investigation of more than 1,500 Ryanair, easyJet and Wizz Air flights" and an "easyJet advertised £5.99 fee" 20. These specific investigations and figures are completely absent from the independent sources in the notebook.
- **Miss-It-Not (ChatGPT Candidate 2)**
- **Pain-Point Evidence Strength: Strong** (general pain point) / **Unsupported** (specific citations) 5, 19.
- **Supporting Source Citations:** *Forensic Field Analysis* (Finding 2) 26, 27, *KAYAK Self-Transfer Guide* 28, 29, *BEUC Factsheet* (No-show policies) 30, 31, *ITV News* (No-show policies) 32, 33.
- **Independent of Proposing Model?** Yes.
- **Audit Gaps & Contradictions:** The report cites specific "NerdWallet" and "Consumer Reports" guidelines for separate-ticket missed connections 5, 34, neither of which exist in the independent sources of the notebook.
- **Senior Fare Copilot (ChatGPT Candidate 3)**
- **Pain-Point Evidence Strength: Unsupported** 35.
- **Supporting Source Citations:** None. General "cognitive overload" is mentioned in *TechRxiv* 36, but there are no selected sources that validate older budget travelers' specific barriers, digital tool dread, or AARP studies.
- **Independent of Proposing Model?** No.
- **Audit Gaps & Contradictions:** Cites "AARP 2026 travel trends" 35, "AARP travel-technology research" 35, and a "2025/2026 smart-tourism digital divide study" 37. None of these documents are present in the independent notebook sources.
- **GateMate (Kimi Candidate 1)**
- **Pain-Point Evidence Strength: Unsupported** 38.
- **Supporting Source Citations:** None. The selected sources do not support senior-specific airport navigation anxiety, digital terminal aversion, or ONS/ACI Europe surveys.
- **Independent of Proposing Model?** No.
- **Audit Gaps & Contradictions:** Cites ONS 2020 internet usage, ACI EUROPE passenger survey on digitalization, and AARP 2024 technology reports 38, all of which are external links and not represented in the independent selected sources.
- **ClaimMate (Kimi Candidate 2)**
- **Pain-Point Evidence Strength: Strong** (general pain point) / **Unsupported** (specific citations) 19, 39.
- **Supporting Source Citations:** *Joint Call on EU Passenger Rights* (BEUC) 40, 41, *Forensic Field Analysis* (Finding 4) 42, 43, *AirHelp* delay rights guide 44, 45.

- **Independent of Proposing Model?** Yes. The BEUC joint call and Forensic Field Analysis are independent.
- **Audit Gaps & Contradictions:** Cites an "AirHelp 2025 global survey" and "European Pulse" stats 39, neither of which are in the independent notebook sources (the AirHelp source in the notebook 44 does not contain this survey). However, the general claims that only 38% of eligible passengers successfully receive compensation and that airlines default to vouchers or deny claims are fully supported by independent sources 40, 43.
- **StitchCheck (Kimi Candidate 3)**
- **Pain-Point Evidence Strength: Strong** 19, 46.
- **Supporting Source Citations:** *Forensic Field Analysis* (Finding 2) 26, 27, *KAYAK Self-Transfer Flights Guide* 28, 29, *BEUC Factsheet on no-shows* 30, 31, *ITV News* no-show ticket cancellations 32, 33.
- **Independent of Proposing Model?** Yes. These are independent of the proposing model's report.
- **Audit Gaps & Contradictions: None.** Every single statistic cited—including the 55% missed connection worry and 21% luggage anxiety from the OAG survey 29, the 30-minute delay causing 40% miss rate from the European Transport Research Review 29, the "no-show" clause cancellations 32, and the €770 average detriment 31—is **fully verified and matched** in the notebook's independent original sources.
- **TrueFare (Z.ai Candidate 1)**
- **Pain-Point Evidence Strength: Strong** (general pain point) / **Unsupported** (specific timeline) 15, 19.
- **Supporting Source Citations:** *BEUC Facts* on baggage 23, *Forensic Analysis* Finding 1 21, 22.
- **Independent of Proposing Model?** Yes.
- **Audit Gaps & Contradictions:** Cites a "Fifth Circuit vacatur of the upfront fee disclosure rule in July 2025" and specific Reuters articles 15. Neither the Fifth Circuit vacatur nor these articles are present in the independent sources of the notebook.
- **ConnectionGuard (Z.ai Candidate 2)**
- **Pain-Point Evidence Strength: Strong** 16, 19.
- **Supporting Source Citations:** *Forensic Field Analysis* (Finding 2) 26, 27, *KAYAK Self-Transfer Guide* 28, 29, *BEUC Factsheet* 30, 31, *ITV News* 32.
- **Independent of Proposing Model?** Yes.
- **Audit Gaps & Contradictions:** Mentions "The Points Guy self-transfer risk explainer" 16. While there is a Points Guy source in the notebook, it is about refunds for canceled flights 47, not self-transfers.
- **DoorFare (Z.ai Candidate 3)**
- **Pain-Point Evidence Strength: Weak / Unsupported** 17, 36.
- **Supporting Source Citations:** None. While general "drip pricing" is mentioned in DOT sources, there is no independent evidence in the selected sources verifying secondary airport transfer distances, shuttles, or travel costs for Frankfurt-Hahn, Beauvais-Tillé, or Girona.
- **Independent of Proposing Model?** No.

- **Audit Gaps & Contradictions:** Cites Wikipedia pages and official airport sites for Frankfurt-Hahn, Beauvais, and Girona 17, none of which are in the independent selected sources of the notebook.

3. Hard-Gate Decision

We evaluated each candidate against the eight non-negotiable eligibility gates. The three official tracks for the Singapore Build with Gemini Hackathon on 22 August 2026 are **Best Use of Gemma**, **Best Elderly Hack**, and **Most Creative Gemini Hack** 9:

Candidate Idea, Theme Fit (Gate 1), Gemini Need (Gate 2), Atlas Need (Gate 3), Nosana Need (Gate 4), 1-Day Feas. (Gate 5), Safety/Demo (Gate 6-8), Overall Eligibility

1. TrueFare Lens, Pass, Pass, Pass, Pass, Pass, Pass, Eligible (Flawed)
2. Miss-It-Not, Pass, Pass, Pass, Pass, Pass, Pass, Eligible (Flawed)
3. Senior Fare Copilot, Pass, Pass, Pass, Pass, Pass, Pass, Fail, Rejected
4. GateMate, Pass, Pass, Pass, Pass, Pass, Pass, Fail, Rejected
5. ClaimMate, Pass, Pass, Pass, Pass, Pass, Pass, Pass, Eligible
6. StitchCheck, Pass, Pass, Pass, Pass, Pass, Pass, Pass, Eligible (Excellent)
7. TrueFare (Z.ai), Fail, Pass, Pass, Pass, Pass, Pass, Rejected
8. ConnectionGuard (Z.ai), Fail, Pass, Pass, Pass, Pass, Pass, Rejected
9. DoorFare (Z.ai), Fail, Pass, Pass, Pass, Pass, Pass, Fail, Rejected

One-Sentence Cited Reasons for Fail / Rejected Decisions:

- **Senior Fare Copilot: Rejected** because the underlying pain point regarding older travelers' digital tool dread and AARP studies is completely unsupported by credible independent evidence in the selected sources (fails Gate 8) 35.
- **GateMate: Rejected** because elder-specific day-of-travel anxiety and digital desk aversion are entirely unsupported by independent selected sources in the notebook (fails Gate 8) 38.
- **TrueFare (Z.ai): Rejected** because its claimed theme "Everyday AI" is not an official Singapore Build with Gemini Hackathon track (fails Gate 1) 9, 15.
- **ConnectionGuard (Z.ai): Rejected** because its claimed theme "The Next Frontier" is not an official Singapore Build with Gemini Hackathon track (fails Gate 1) 9, 16.
- **DoorFare (Z.ai): Rejected** because its claimed theme "Social Good" is not an official hackathon track (fails Gate 1) 9, 18, and transit cost data for specific secondary airports is completely unsupported by the notebook's independent sources (fails Gate 8) 17.

4. Score Eligible Ideas Only

We scored the four eligible candidates on a 1–5 scale across nine distinct criteria:

Evaluation Criterion, TrueFare Lens (ChatGPT 1), Miss-It-Not (ChatGPT 2), ClaimMate (Kimi 2), StitchCheck (Kimi 3)

1. Pain Severity & Evidence Quality, 3/5, 3/5, 4/5, 5/5
2. User-Value Clarity, 4/5, 4/5, 3/5, 5/5
3. Originality & Theme Fit, 3/5, 4/5, 4/5, 5/5
4. Gemini Integration Depth, 3/5, 3/5, 4/5, 5/5

5. Atlas Sandbox Integration Depth,3/5,3/5,3/5,5/5
6. Nosana Integration Depth,3/5,4/5,3/5,5/5
7. One-day Build Feasibility,4/5,4/5,4/5,5/5
8. Demo Clarity & Reliability,3/5,4/5,3/5,5/5
9. Safety & Trust Design,4/5,4/5,2/5,5/5
- TOTAL SCORE (out of 45),30/45,33/45,30/45,45/45

Score Justification Highlights:

- **StitchCheck (45/45):** Outstanding. Every statistic and study cited is 100% verified by the independent sources (KAYAK 29, BEUC 31, and ITV 32). Gemini, Atlas, and Nosana all have essential, heavy-lifting, and visible roles. The design enforces synthetic-only P0 data and explicit human verification for write actions, making it exceptionally safe 48, 49.
- **Miss-It-Not (33/45):** A strong creative concept, but penalized due to the inclusion of fabricated/missing "NerdWallet" and "Consumer Reports" source references in its report 5, and the lack of Atlas write actions in its vertical slice 50.
- **TrueFare Lens (30/45):** Penalized because its specific Which? December 2025 investigation claims are fabricated 20, and it features no write actions or priced rebooking flows in its P0 slice 51.
- **ClaimMate (30/45):** Severely penalized on Safety & Trust (2/5). It claims to be a "privacy-first flight-compensation copilot" running Gemma to protect user IDs and travel documents 11, 12. However, deploying this workload onto Nosana's decentralized, public GPU network means sending sensitive personal documents (disruption emails, booking receipts, IDs) to untrusted public nodes, introducing massive security and privacy risks.

5. Winning Recommendation

We are thrilled to recommend **StitchCheck** (Origin: **research-kimi.md**) as the unanimous winning idea 13. It passes every hard gate flawlessly, stands on an unshakeable, verified evidence base, and leverages the entire tech stack with unparalleled elegance.

1. **Idea Name & Source Model: StitchCheck** by Kimi 13.
2. **Selected Hackathon Theme: Most Creative Gemini Hack** (Official Track) 9, 14.
3. **One-Sentence Pitch:** Screenshot your cheap separate tickets before paying, and StitchCheck instantly builds a narrated "failure-cascade timeline" to expose your real missed-connection risks and book a safer alternative in the sandbox 13.
4. **Why the User Pain is Real (Independent Evidence):**
5. **Extreme vulnerability of self-connections:** A minor flight delay of just 30 minutes causes **40%** of self-connecting passengers to miss their onward segments 27, 29.
6. **No legal fallback:** Point-to-point contracts on separate tickets are completely exempt from through-ticket rebooking and care protections under both EU261 and US DOT rules 26, 52, 53.
7. **Punitive "No-Show" cascades:** Missing the first leg frequently triggers "no-show" clauses, allowing low-cost carriers to cancel the traveler's entire return itinerary without refunds 27, 30.

8. **High financial harm:** Stranded travelers face average financial losses of **€770** per disrupted booking when airlines deny care and rerouting 27, 31.
9. **Gemini Role & Visible Proof:**
10. *Role:* Gemini (Flash-class) performs multimodal extraction on screenshots of separate checkouts/tickets to build a structured itinerary JSON, then maps these flights against the destination airport's minimum connection layout to generate a spoken "failure-cascade narrative" 14, 48.
11. *Visible proof in demo:* The UI displays the original ticket screenshots side-by-side with a rendered gate-to-gate transit schedule and the Gemini-narrated timeline of what happens minute-by-minute if Leg 1 is delayed 48.
12. **Atlas Sandbox Role & Visible Proof:**
13. *Role:* Atlas searches real flight shopping tables to retrieve safer, alliance-linked, or protected connecting alternatives, and handles the mock-booking checkout 48.
14. *Visible proof in demo:* A priced table of Atlas alternative itineraries is displayed (e.g., the risky separate-ticket connection vs. a safer protected connection), and a sandbox booking confirmation screen appears when the traveler decides to secure the flight 48.
15. **Nosana Role & Visible Proof:**
16. *Role:* Executes high-throughput batch Monte Carlo reliability simulations on route-pair delay data to serve a risk-score API. This is a perfect workload for decentralized GPUs as it involves heavy numerical calculations with zero privacy risks (no PII involved) 54.
17. *Visible proof in demo:* A live status panel displays the Nosana Job ID, GPU node details, and the finalized connection risk score rendered in the UI 54.
18. **P0 Vertical Slice (Single User Journey):**
19. *One User:* Synthetic traveler profile ("Demo Traveler / demo@example.com") 55.
20. *One Trigger:* Traveler uploads screenshots of two unbooked, tight separate-ticket checkouts 54.
21. *One Decision:* StitchCheck reveals the 75-minute connection has a **40%** risk score 29, 54, and presents a safer protected route for only €14 more 16. The user must decide to risk it ("Keep") or choose safety ("Switch") 54.
22. *One Outcome:* The user clicks "Switch" to buy the safer option, triggering a human-confirmed sandbox flight booking receipt 48, 54.
23. **Explicit Non-Goals:**
24. No integration with real payment cards or production checkout gateways.
25. No live, dynamic real-time prediction of weather/delays (uses a static historical dataset) 49.
26. No actual production bookings or passenger seat reservations on live flights.
27. No storage or processing of real personal identifying information (PII).
28. **Three Highest-Risk Assumptions to Smoke-Test First:**
29. *Screenshot OCR Robustness:* Gemini's ability to accurately extract exact departure/arrival times, airport codes, and carrier names across highly fragmented layout variations of various airline checkout pages 49.
30. *Historical Data Coverage:* That the historical delay dataset used by the Nosana batch job has sufficient coverage for secondary LCC airports to produce meaningful risk heuristics 49.

31. *Atlas Sandbox API Latency:* That the Atlas Flight Booking Sandbox API responds reliably and quickly under hackathon network constraints to return alternative priced offers 49.

 **Next Step:** If this selection panel's evaluation aligns with your vision, I can programmatically compile this evaluation into a polished presentation file or detailed PDF report that you can present directly to the hackathon organizers. Would you like me to do that?